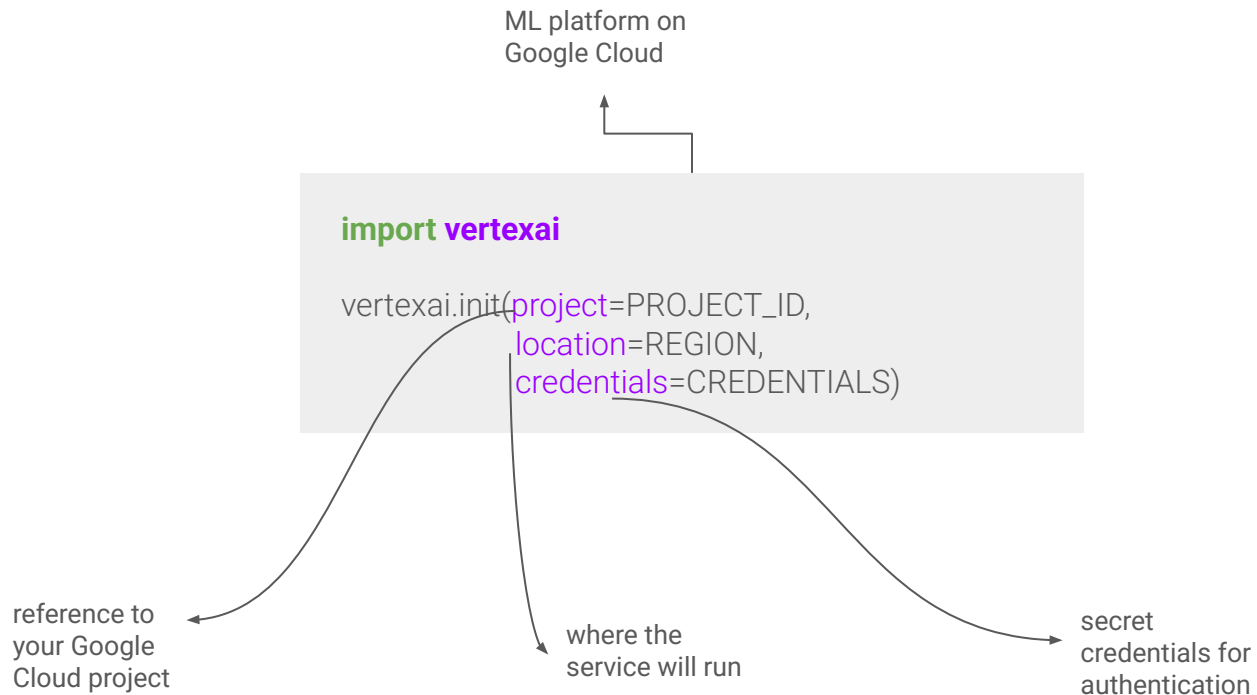


Understanding and Applying Text Embeddings

Set up Google Cloud



Specify a model, Get an Embedding

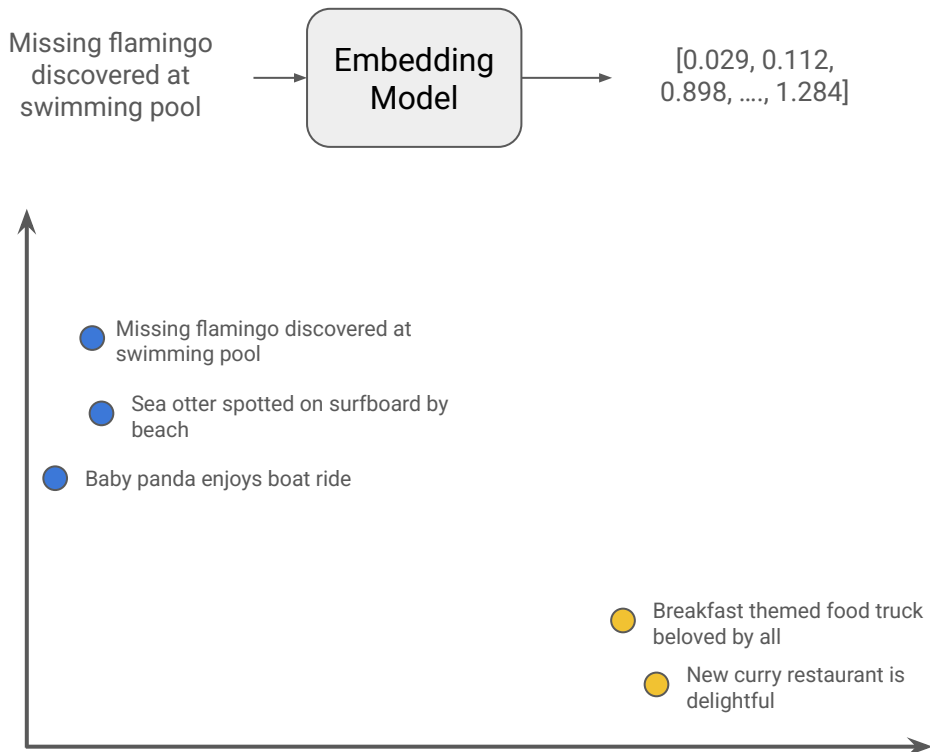
```
from vertexai.language_models import TextEmbeddingModel
```

```
embedding_model = TextEmbeddingModel.from_pretrained(  
    "text-embedding-005")
```

```
embedding = embedding_model.get_embeddings(  
    ["What is the meaning of life?"])
```

What is an embedding

A way to represent data as points in space where the locations are semantically meaningful



How are sentence embeddings computed

- **Simple Method:**
 - Embed each word separately, and take a sum or mean of all the word embeddings
- **Modern Embeddings:**
 - Use a transformer neural network to compute a context-aware representation of each word, then take an average of the context-aware representations
 - Compute embeddings for each token (e.g., sub-word) rather than word. Enables algorithm to work even for novel words and misspelt words (“Life the universe and everything”).
- **Training the transformer network (contrastive learning)**
 - Given a dataset of pairs of “similar” sentences, tune neural network to move similar sentences embeddings together and dissimilar sentences’ embeddings apart.

Specify a mode, Generate Text

```
from vertexai.generative_models import GenerativeModel
```

```
model = GenerativeModel("gemini-2.5-flash")
```

```
prompt = "Recommend me a programming activity to improve my skills."  
response = model.generate_content(contents=prompt)
```

Decoding Strategies

Text Input

The garden was full of beautiful

Probabilities over tokens

[flowers (0.5),
trees (0.23),
herbs (0.05),
.....,
bugs (0.03)]



Which one should you
choose?

Decoding Strategies

Text Input

The garden was full of beautiful flowers

Probabilities over tokens

[flowers (0.5),
trees (0.23),
herbs (0.05),
.....,
bugs (0.03)]

Greedy Decoding

The one with the highest probability

Decoding Strategies

Text Input

The garden was full of beautiful bugs

Probabilities over tokens

[flowers (0.5),
trees (0.23),
herbs (0.05),
.....,
bugs (0.03)]




Random sample

Use the probabilities to sample a
random token

Temperature

Text Input

The garden was full of beautiful



Temperature controls the randomness

Temperature

Word	Logits	Softmax	Softmax with temperature 0.1
flowers	20	0.881	1.000
trees	18	0.119	0.000
herbs	5	0.000	0.000
bugs	2	0.000	0.000

Softmax

$$\text{Softmax}_{\theta}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$


Softmax with temperature θ

$$\text{Softmax}_{\theta}(z_i) = \frac{e^{\frac{z_i}{\theta}}}{\sum_{j=1}^n e^{\frac{z_j}{\theta}}}$$

Temperature

Text Input

The garden was full of beautiful




Temperature of 0:
Deterministic choice
Not creative but reliable

Temperature

Text Input


The garden was full of beautiful



As temperature increases so does randomness

Creative but sometimes unreasonable

Temperature



Top K

Text Input

The garden was full of beautiful

Probabilities over tokens

[flowers (0.5),
trees (0.23),
herbs (0.05),
.....,
bugs (0.03)]

Sample from tokens with top k
probabilities


Top K

Text Input

The garden was full of beautiful

Probabilities over tokens

[flowers (0.5),
trees (0.23),
herbs (0.05),
.....,
bugs (0.03)]




$k = 2$

Sample from tokens with top k probabilities

Top K

Text Input


It is raining cats and _____



Top K

Text Input

It is raining cats and hail



Text Input

The garden was full of beautiful

Probabilities over tokens

[flowers (0.5),
trees (0.23),
herbs (0.05),
.....,
bugs (0.03)]


Sample from minimum set of
tokens whose cumulative
probability is greater or equal to P

Text Input

The garden was full of beautiful

Probabilities over tokens

[flowers (0.5),
trees (0.23),
herbs (0.05),
.....,
bugs (0.03)]



$P = 0.75$

Sample from minimum set of tokens whose cumulative probability is greater or equal to P

Putting it all together

Text Input

The garden was full of beautiful

[flowers, trees, herbs, plants, _ _ _ , weeds, bugs]

Top K

[flowers, trees, herbs, plants]

Top P

[flowers, trees]

Temperature

trees

Putting it all together

```
from vertexai.generative_models import GenerativeModel, GenerationConfig
```

```
model = GenerativeModel("gemini-2.5-flash")
```

```
config= GenerationConfig(temperature=0.9, top_p=0.7, top_k=20)
```

```
prompt = "Recommend me a programming activity to improve my skills."  
response = model.generate_content(contents=prompt, generation_config=config)
```

Grounding LLMs

Out-of-the-box LLM aren't connected to real world

```
from vertexai.generative_models import GenerativeModel, GenerationConfig
```

```
prompt = "How to concat dataframes in pandas"  
response = model.generate_content(contents=prompt, generation_config=config)
```

Response would depend on the knowledge cut-off

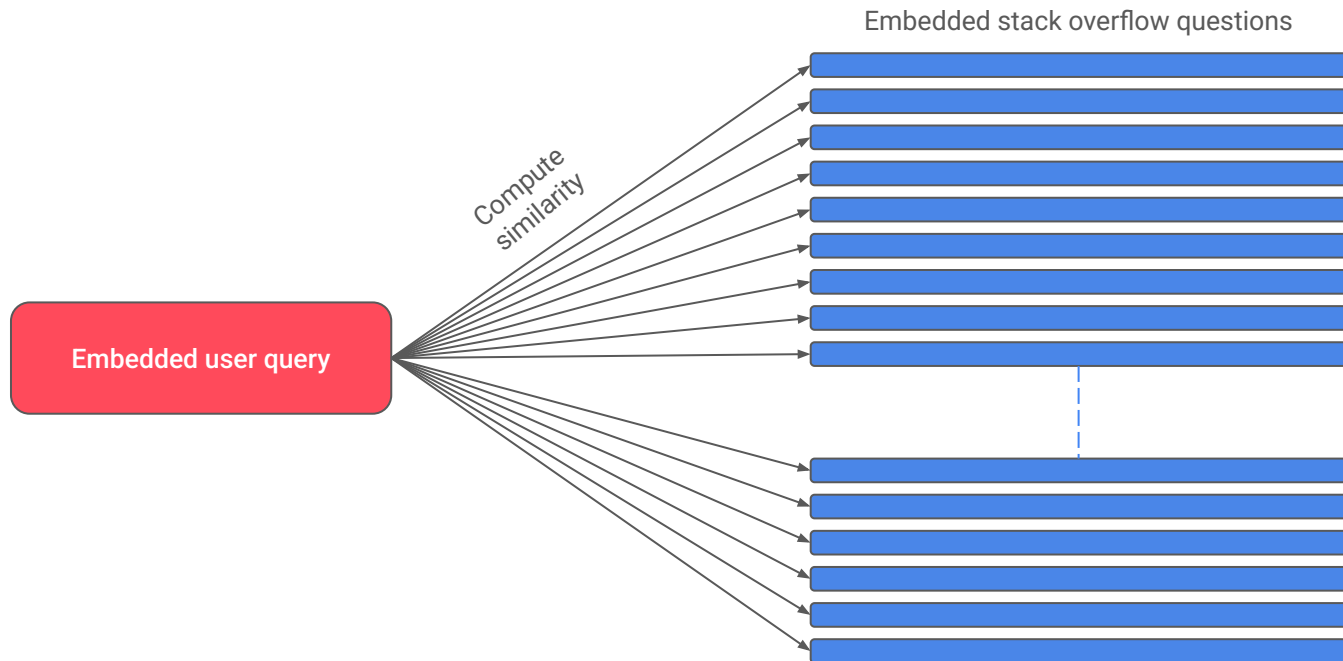
Grounding LLMs:

- Access information outside of training data
- Integrate with existing IT systems, databases, and business data
- Mitigate risk of hallucinations


Q&A over Stack Overflow data

Question: Combine 3 dataframes in pandas. I have a dataframe like this

Accepted Answer: You can use `pd.concat` here. Just define and pass all 3 dataframes to . . .



Measuring Similarity using Embeddings




Euclidean Distance (L2 Distance): Distance between ends of vectors

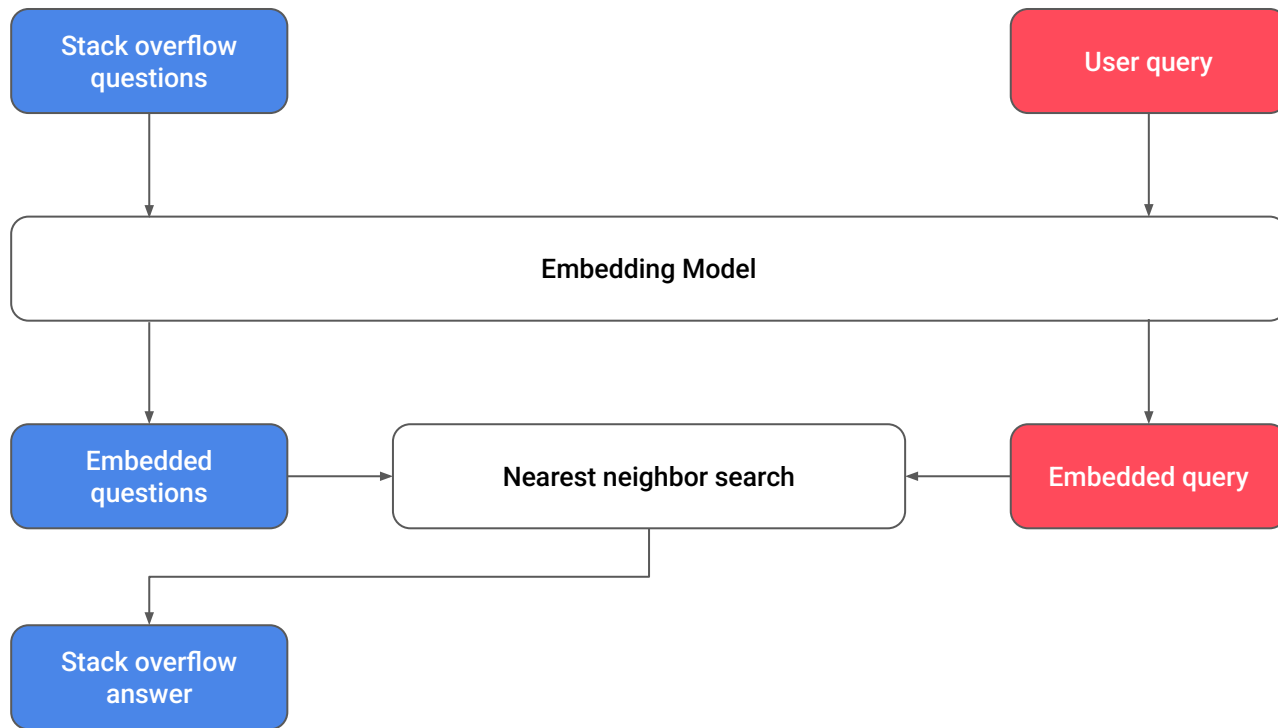
Cosine Similarity: Cosine of the angle between vectors

Dot Product: Cosine multiplied by lengths of both vectors

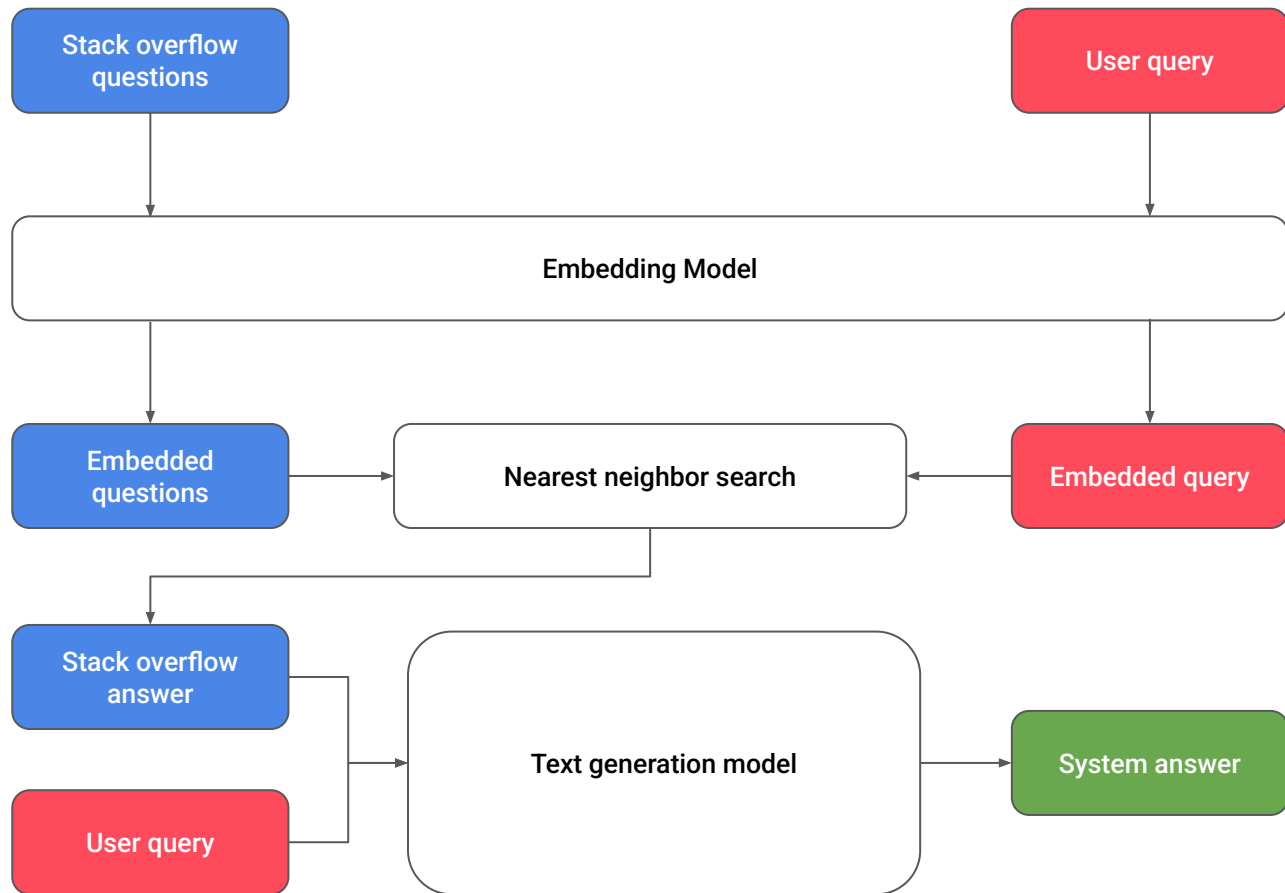
Q&A with Semantic Search



Q&A with Semantic Search



Q&A with Semantic Search



THANK YOU